

SEO Content & E-E-A-T Audit

mangoaiusa.com — homepage, with /about and /blog cross-checks

40/100 CONTENT QUALITY	34/100 AI CITATION READINESS	2,118 WORDS RENDERED	0 JSON-LD BLOCKS
---------------------------	---------------------------------	-------------------------	---------------------

Audited 24 August 2026 · Method: raw HTML fetch with no JavaScript execution, matching what crawlers and LLM fetchers see. Scores are this framework's heuristics, weighted to Google's stated hierarchy (Trust 30 / Expertise 25 / Authoritativeness 25 / Experience 20). They are not Google-internal signals; no third-party tool has access to those.

Google's Who / How / Why test

QUESTION	VERDICT
Who created it?	No named author or human identity anywhere on the site beyond "Mango AI, Inc." The About page says "we are a small team" and names nobody. Blog bylines read "The Mango team."
How was it created?	No process disclosure, no original data, no research. Every number on the page — 4.9/5, 500+ teams, +41% weekend revenue — is asserted without source or method.
Why does it exist?	Genuinely helpful framing (the "capacity problem at the exact moment a customer is ready to buy" is a real insight, well written), but it is pure conversion copy with no informational payload.

E-E-A-T breakdown

FACTOR	SCORE	KEY SIGNALS
Experience	6 / 20	No case studies, no before/after data, no screenshots of real accounts. The four-step call walkthrough uses a fabricated scenario — "Dana Reyes, 128 Cedar St., +1 (555) 0198." A 555 number signals <i>illustration</i> , not experience. The ROI calculator is the one genuine first-party asset.
Expertise	12 / 25	Strong trade fluency: burst pipes, water heater failure, sewer backups, emergency-versus-routine triage, "the phone rings while everyone is already on a job." That is someone who knows the vertical. Undercut by zero author credentials and three roughly one-minute blog posts.
Authoritativeness	7 / 25	Zero external citations or outbound links. Customer logos (AquaPlumb, AirPro HVAC, VoltElectric, SafeGuard, BlueLine) are unlinked and do not read as verifiable businesses. No press, no G2 or Capterra, no reviews platform.
Trustworthiness	15 / 30	HTTPS, plus live privacy, terms and security pages. Honest hedging throughout ("Estimate only", "Coming soon: direct sync"), and the line "Every promise below is ours to keep — not a rating someone else gave us" is unusually candid and earns real credit. Against that: no contact page (/contact returns 404), and no address, phone or email anywhere. For a product whose entire pitch is "we answer your phone," having no reachable phone number is a conspicuous gap.

AI citation readiness

SIGNAL	STATUS
Structured data	Zero JSON-LD on any page. No Organization, SoftwareApplication, Product or Offer schema.
Heading hierarchy	Clean H1 → H2 → H3, descriptive and scannable. The strongest signal on the page.
FAQ extractability	Only 1 of 6 FAQ answers is present in the HTML. The other five render client-side on accordion open, invisible to crawlers and LLM fetchers.
Key statistics	Animated counters ship as zeros. The outcomes band literally serves 0% more jobs booked, 0s average time to book, \$0 saved every month, 0+ home service teams. Anything reading raw HTML sees a company claiming zero results.
Comparative tables	Pricing tiers are real text — Starter \$79 + \$0.50/call, Pro \$239 with 500 calls then \$0.40, Enterprise custom. Highly quotable.
Attribution	Testimonials carry name, title and company, but no links and no verification path.
Canonical identity	Points at a different domain. og:url is https://mango.ai, and sitemap.xml lists every URL on mango.ai rather than mangoaiusa.com. No <link rel="canonical"> at all.

Issues found

CRITICAL

1. **CRIT Domain identity conflict.** The sitemap served at `mangoaiusa.com/sitemap.xml` declares eight or more URLs, all on `mango.ai`. Combined with `og:url=https://mango.ai` and no canonical tag, the site is telling search engines its real home is a different domain. Either that domain is the canonical property — in which case this one should redirect — or the metadata is stale.
2. **CRIT robots.txt returns 404.** No crawl directives and no sitemap declaration.
3. **CRIT Counter stats render as 0** in server HTML. The four headline proof points are self-defeating for crawlers and AI fetchers alike.
4. **CRIT Five of six FAQ answers absent from HTML.** The most citable content block on the page is largely unreadable to machines.
5. **CRIT No structured data sitewide** — homepage, /about and /blog all return zero JSON-LD blocks.

HIGH

6. **HIGH** No contact page (/contact 404), no NAP, no support email or phone.
7. **HIGH** No named humans anywhere — no leadership, no author bios, no LinkedIn.
8. **HIGH** Unattributed proof claims: "4.9/5", "Trusted by 500+ home service teams", "+41% weekend revenue", "6 hrs/day saved".
9. **HIGH** Blog is thin — three posts, all "1 min read", all bylined "The Mango team", no schema.
10. **HIGH** Pricing inconsistency: the ROI calculator says "\$99/mo + 50¢ per call" while the pricing table says Starter is "\$79/mo + \$0.50 per call."

MEDIUM

11. **MED** Only two `<img>` tags on the page, both with alt text. Nearly all visuals are CSS/SVG composites, so there is no image-search surface and nothing for AI multimodal retrieval.
12. **MED** The H1 is split by animation — "Never miss another" is in the HTML, "service call. Ever." arrives via JavaScript. The crawled H1 is a fragment.
13. **MED** Every internal link is navigation or footer. No contextual in-body links to /platform, /industries or blog posts.
14. **MED** `<meta name="keywords">` is present and ignored by every engine. Harmless in itself, but it signals the SEO setup was never audited.

Recommendations

THIS WEEK

1. Fix the sitemap to emit `mangoaiusa.com` URLs, add `<link rel="canonical">` to every page, and correct `og:url`. If `mango.ai` is the real property, 301 this domain instead.
2. Add `robots.txt` with Sitemap: `https://mangoaiusa.com/sitemap.xml`.

3. Server-render the counter final values and hydrate the animation from them — the `<span data-target="38">38</span>` pattern. Never ship `0` as the base.
4. Render all six FAQ answers into the DOM. Hide them with CSS, not conditional mounting.
5. Build `/contact` with a real address, phone and support email.

THIS MONTH

6. Add JSON-LD: `Organization` (with `sameAs` and `contactPoint`), `SoftwareApplication` plus `offer` for the three tiers, and `BreadcrumbList`. Skip `FAQPage` — it no longer produces Google rich results — but the visible answer text still feeds AI extraction.
7. Put names and faces on it: a founder or ops lead on `/about`, real bylines on blog posts, `Person` schema for each.
8. Source or retire the numbers. One named customer with a verifiable outcome — "Apex Plumbing, Tacoma WA: 47 after-hours calls captured in month one" — is worth more than four round unattributed stats, and is the single highest-leverage fix for both Experience and Authoritativeness.
9. Reconcile the \$99 versus \$79 Starter price.
10. Turn the blog into real coverage. Expand "What a missed call actually costs" into a 1,500-word piece built on your own aggregate call data. That is the asset AI engines will actually cite, and you are already sitting on the dataset.

Worth keeping. The heading structure, the trade-specific vocabulary, the pricing table and the "not a rating someone else gave us" honesty line are all genuinely good. The gap here is not writing quality — it is that almost none of the trust and machine-readability infrastructure exists yet.

On word count and readability. Google has confirmed word count is not a ranking factor, and John Mueller has confirmed basic readability scores are not used for ranking. The 2,118-word homepage clears any reasonable coverage floor; depth of evidence, not length, is what is missing.